

Big data science Day 3

F. Legger - INFN Torino

<https://github.com/leggerf/MLCourse-2022>



We learned

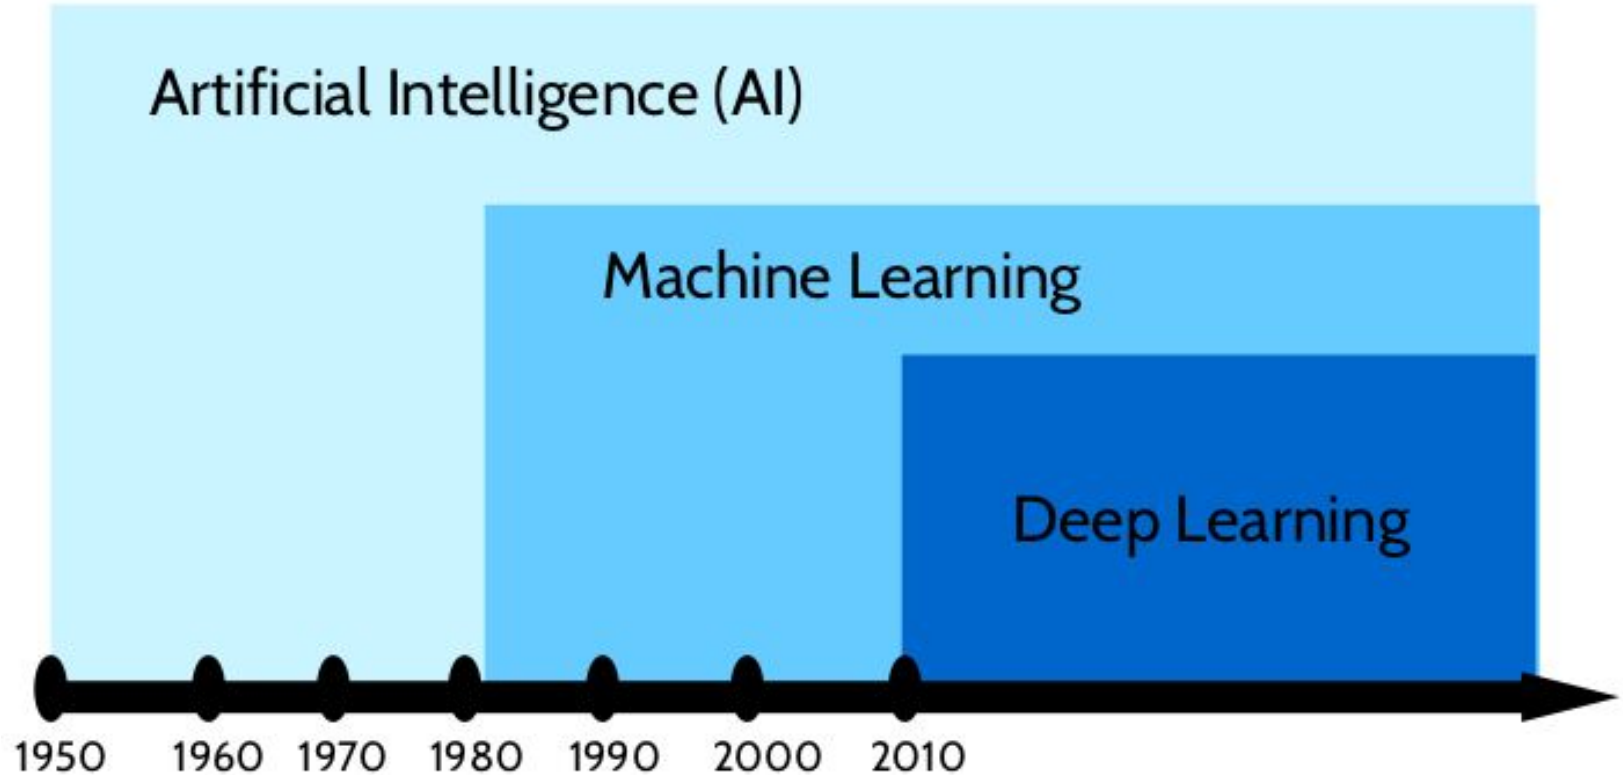
- Big data
- Analytics
- Machine learning

Today

- Deep learning
- Parallelisation
- Heterogeneous architectures
- Future directions




Deep Learning is a subfield of ML concerned with algorithms inspired by the structure and function of the brain called artificial neural networks [Jason Brownlee]



Machine translation

Real-time translation into Mandarin Chinese (2012)




Strategy games

DeepMind beats Go world champion (2017)



Creativity




Self driving cars



Visual recognition



Short recap: Neuron



- each **input** x is multiplied by a **weight** w

$$x_1 \rightarrow x_1 * w_1$$



$$x_2 \rightarrow x_2 * w_2$$

- all the weighted inputs are added together with a **bias** b

$$(x_1 * w_1) + (x_2 * w_2) + b$$




- the sum is passed through an **activation function** f

$$y = f(x_1 * w_1 + x_2 * w_2 + b)$$



Neural network

- Combining more neurons
- A **hidden layer** is any layer between the input (first) layer and output (last) layer
 - There can be multiple hidden layers
- **Feedforward**: process of passing inputs forward to get an output




This network has:

- one **input** layer with 2 inputs
- one **hidden** layer with 2 neurons
- one **output** layer with 1 neuron

Deep Learning


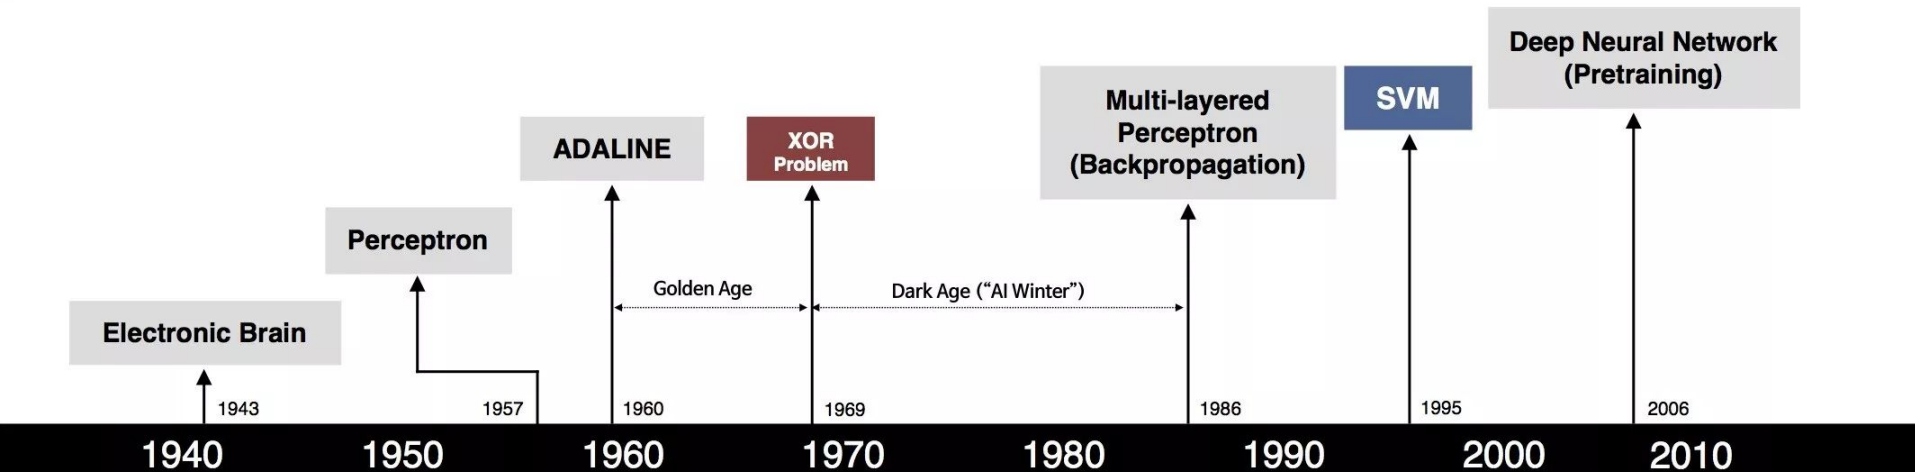
- Neural network with several layers
 - Deep vs shallow
- A family of parametric models which learn non-linear hierarchical representations:



$$a_L(\mathbf{x}; \Theta) = h_L(h_{L-1}(\dots(h_1(\mathbf{x}, \theta_1), \theta_{L-1}), \theta_L)$$

input parameters of the network non-linear activation function parameters of layer L

Brief history of neural networks



S. McCulloch - W. Pitts



F. Rosenblatt



B. Widrow - M. Hoff



M. Minsky - S. Papert




D. Rumelhart - G. Hinton - R. Williams




V. Vapnik - C. Cortes




G. Hinton - S. Ruslan




- Adjustable Weights
- Weights are not Learned




- Learnable Weights and Threshold




- XOR Problem



- Solution to nonlinearly separable problems
- Big computation, local optima and overfitting




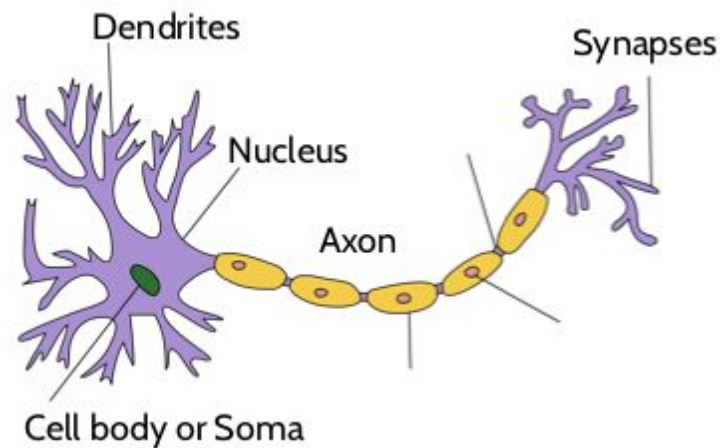
- Limitations of learning prior knowledge
- Kernel function: Human Intervention



- Hierarchical feature Learning


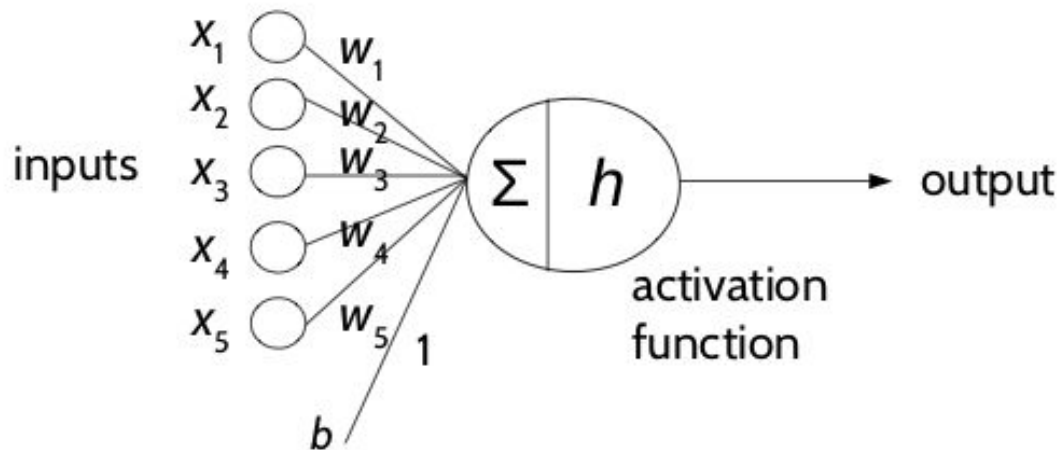
1943 – McCulloch & Pitts Model

- Early model of artificial neuron
- Generates a binary output
- The weights values are fixed



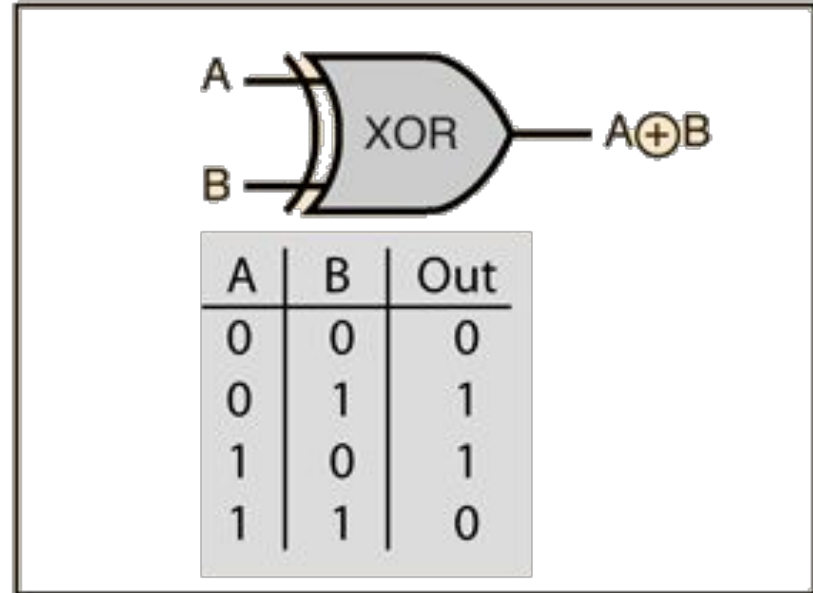
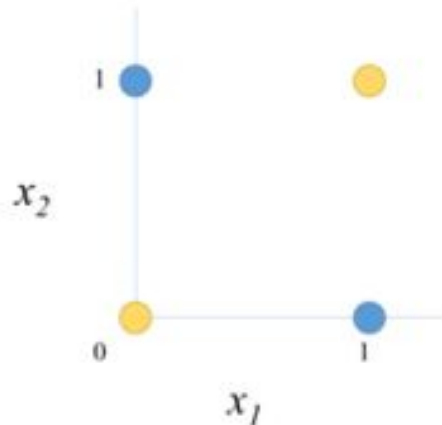
1958 – Perceptron by Rosemblatt

- Perceptron as a machine for linear classification
- Main idea: Learn the weights and consider bias.
 - One weight per input
 - Multiply weights with respective inputs and add bias
 - If result larger than **threshold** return 1, otherwise 0




First NN winter

- 1970- Minsky. The XOR cannot be solved by perceptrons.
- Neural models cannot be applied to complex tasks.



Multi-layer Feed Forward Neural Network

- 1980s. Multi-layer Perceptrons (MLP) can solve XOR.
- ML Feed Forward Neural Networks:
 - Densely connect artificial neurons to realize compositions of non-linear functions
 - The information is propagated from the inputs to the outputs
 - The input data are usually n -dimensional feature vectors
 - Tasks: Classification, Regression



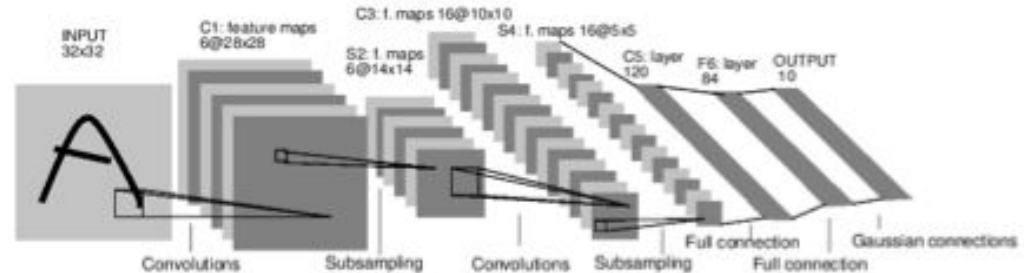
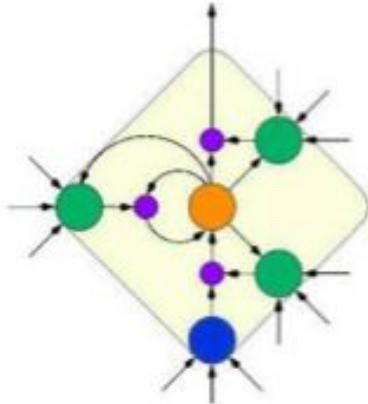
How to train it?

- Rosenblatt algorithm* not applicable, as it expects to know the desired target
 - For hidden layers we cannot know the desired target
- Learning MLP for complicated functions can be solved with **Back propagation (1980)**
 - efficient algorithm for complex NN which processes large training sets

* Remember? Rosenblatt developed a method to train a single neuron

1990s - CNN and LSTM

- Important advances in the field:
 - Backpropagation
 - Recurrent Long-Short Term Memory Networks (Schmidhuber, 1997)
 - Convolutional Neural Networks - LeNet: OCR solved before 2000s (LeCun, 1998).



OCR: Optical character recognition

Convolutional Neural Networks (CNN)

- **Convolutional** layer: two functions produce a third that describes how the shape of one is changed by the other
- **pooling** layer: reduce dimensionality

Source layer

5	2	6	8	2	0	1	2
4	3	4	5	1	9	6	3
3	9	2	4	7	7	6	9
1	3	4	6	8	2	2	1
8	4	6	2	3	1	8	8
5	8	9	0	1	0	2	3
9	2	6	6	3	6	2	1
9	8	8	2	6	3	4	5

Convolutional kernel

-1	0	1
2	1	2
1	-2	0

Destination layer

		5					

$$\begin{aligned} &(-1 \times 5) + (0 \times 2) + (1 \times 6) + \\ &(2 \times 4) + (1 \times 3) + (2 \times 4) + \\ &(1 \times 3) + (-2 \times 9) + (0 \times 2) = 5 \end{aligned}$$



4	1	5	0
7	8	9	8
3	5	6	5
2	4	1	0

Max pooling

8	9
5	6


Second NN Winter

- NN cannot exploit many layers
 - Overfitting
 - Vanishing gradient (with NN training you need to multiply several small numbers → they become smaller and smaller)
- Lack of processing power (no GPUs)
- Lack of data (no large annotated datasets)
- Kernel Machines (e.g. SVMs) suddenly become very popular◦




ImageNet

A Large-Scale Hierarchical Image Database (2009)

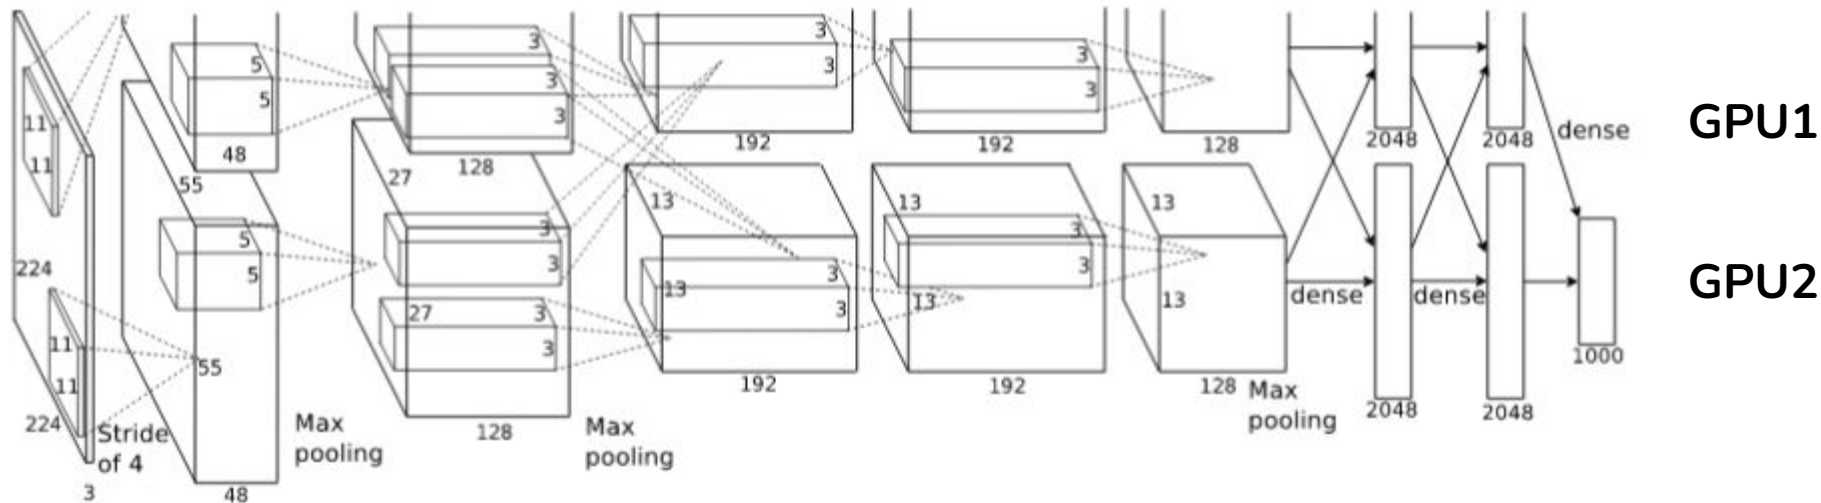


2012 - AlexNet

- Hinton's group implemented a CNN similar to LeNet [LeCun1998] but...
 - Trained on ImageNet (1.4M images, 1K categories)
 - With 2 GPUs
 - Other technical improvements (ReLU, dropout, data augmentation)



AlexNet

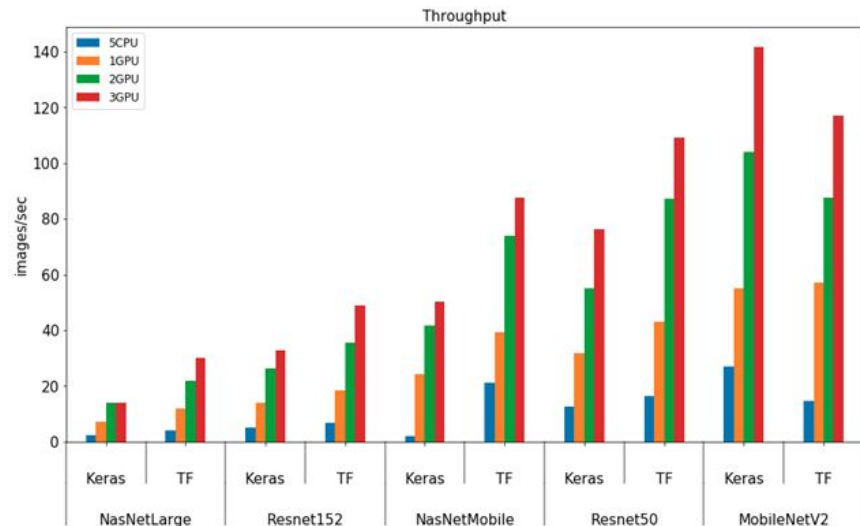


- 60M parameters
- Limited information exchange between GPUs

Why Deep Learning now?

- Three main factors:
 - Better hardware
 - Big data
 - Technical advances:
 - Layer-wise pretraining
 - Optimization (e.g. Adam, batch normalization)
 - Regularization (e.g. dropout)


....



Rectified Linear Units - Activation function (2010)


$$f(x) = \max(0, x)$$

- More efficient gradient propagation: (derivative is 0 or constant)
- More efficient computation: (only comparison, addition and multiplication).
- Sparse activation: e.g. in a randomly initialized networks, only about 50% of hidden units are activated (having a non-zero output)



Regularization - Dropout

- For each instance drop a node (hidden or input) and its connections with probability p and train
- Final net just has all averaged weights (actually scaled by $1-p$)
- As if ensembling 2^n different network substructures



Data augmentation

- Techniques to significantly increase the diversity of data available for training models, without actually collecting new data
- Data augmentation techniques such as cropping, padding, and horizontal flipping are commonly used to train large neural networks




Original



Horizontal Flip



Pad & Crop



Rotate

Training a neural network

Name	Weight (lb)	Height (in)	Gender
Alice	133	65	F
Bob	160	72	M
Charlie	152	70	M
Diana	120	60	F

- Predict gender from weight and height

Feature engineering

- Symmetrize numeric values
- Category -> numbers

Name	Weight (lb)	Height (in)	Gender
Alice	133	65	F
Bob	160	72	M
Charlie	152	70	M
Diana	120	60	F


Name	Weight (minus 135)	Height (minus 66)	Gender
Alice	-2	-1	1
Bob	25	6	0
Charlie	17	4	0
Diana	-15	-6	1



Ingredients

- n : 4, number of samples (Alice, Bob, Charlie, Diana)
- y : variable being predicted (Gender)
- y_{true} : true value of y , y_{pred} : predicted value of y = \bullet
- Loss function **L: MSE**
- Activation function **f**
- Outputs of the hidden layer **h**
- Unknown parameters: weights **w** and biases **b**

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_{\text{true}} - y_{\text{pred}})^2$$



Back propagation

- Training the network == trying to minimize its loss
 - Find weights \mathbf{w} and biases \mathbf{b}
 - $L(w_1, w_2, w_3, w_4, w_5, w_6, b_1, b_2, b_3)$
- Minimization taking partial derivatives (**back propagation**)

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$$

For very simple case: with only Alice in the dataset, $n=1$

The diagram illustrates the backpropagation of gradients through three equations. A central box contains the chain rule formula: $\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial y_{pred}} * \frac{\partial y_{pred}}{\partial h_1} * \frac{\partial h_1}{\partial w_1}$. Three blue arrows point from this box to three separate equations below. The first equation is the loss function $L = (1 - y_{pred})^2$, with the partial derivative $\frac{\partial L}{\partial y_{pred}} = \frac{\partial (1 - y_{pred})^2}{\partial y_{pred}}$ shown to its left. The second equation is the prediction $y_{pred} = o_1 = f(w_5 h_1 + w_6 h_2 + b_3)$, with the partial derivative $\frac{\partial y_{pred}}{\partial h_1} = w_5 * f'(w_5 h_1 + w_6 h_2 + b_3)$ shown to its left. The third equation is the hidden layer output $h_1 = f(w_1 x_1 + w_2 x_2 + b_1)$, with the partial derivative $\frac{\partial h_1}{\partial w_1} = x_1 * f'(w_1 x_1 + w_2 x_2 + b_1)$ shown to its left. Vertical blue lines separate the three equations.


$$L = (1 - y_{pred})^2$$
$$\frac{\partial L}{\partial y_{pred}} = \frac{\partial (1 - y_{pred})^2}{\partial y_{pred}}$$
$$y_{pred} = o_1 = f(w_5 h_1 + w_6 h_2 + b_3)$$
$$\frac{\partial y_{pred}}{\partial h_1} = w_5 * f'(w_5 h_1 + w_6 h_2 + b_3)$$
$$h_1 = f(w_1 x_1 + w_2 x_2 + b_1)$$
$$\frac{\partial h_1}{\partial w_1} = x_1 * f'(w_1 x_1 + w_2 x_2 + b_1)$$

Gradient Descent

- optimization algorithm to find weights and biases to minimize loss
- **update equation:**

$$w_1 \leftarrow w_1 - \eta \frac{\partial L}{\partial w_1}$$

- **η (learning rate)** is a constant that controls how fast we train
- If $\frac{\partial L}{\partial w_1}$ is positive, w_1 will decrease, which makes L decrease.
- If $\frac{\partial L}{\partial w_1}$ is negative, w_1 will increase, which makes L decrease.




Stochastic gradient descent (SDG)

- **Stochastic** -> the parameters are updated using only a single training instance (usually randomly selected) in each iteration
- Use mini-batch **sampl**ed in the dataset for gradient estimate.

$$\Theta^{t+1} = \Theta^t - \frac{\eta_t}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \nabla_{\Theta} \mathcal{L}_i$$


- Sometimes helps to escape from local minima
- Noisy gradients act as regularization
- Variance of gradients increases when batch size decreases
- Not clear how many sample per batch

— Batch gradient descent
— Mini-batch gradient Descent
— Stochastic gradient descent



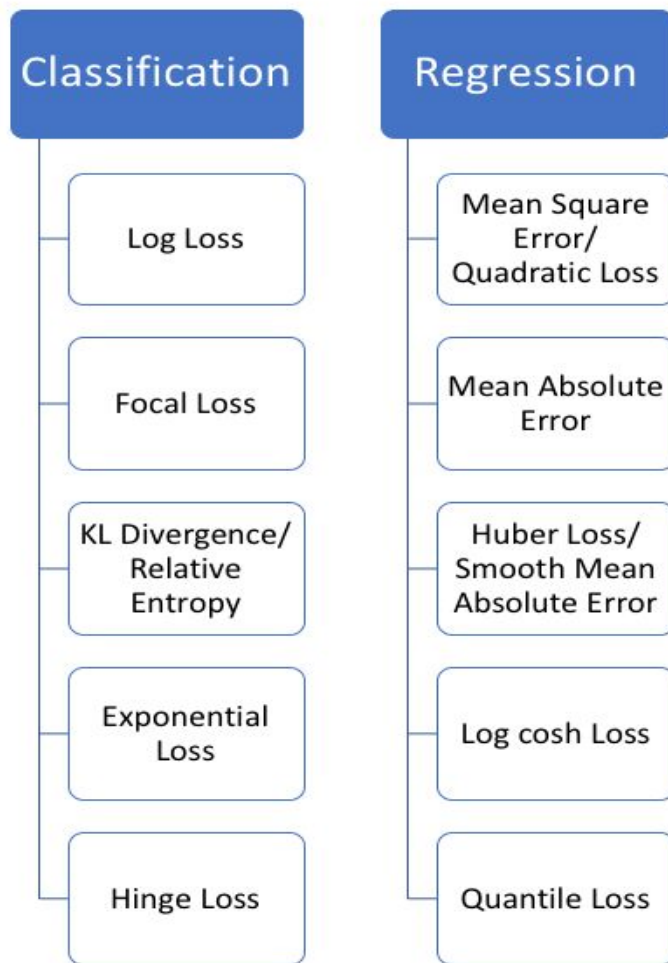
Training the network

- Choose one sample from our dataset
 - This is what makes it stochastic gradient descent - only operate on one sample at a time
- Calculate all the partial derivatives of loss with respect to weights or biases
- Use the update equation to update each weight and bias
- Iterate



Loss functions

- <https://heartbeat.fritz.ai/5-regression-loss-functions-all-machine-learners-should-know-4fb140e9d4b0>
- https://www.wikiwand.com/en/Loss_functions_for_classification



Regularization


- One of the major aspects of training the model is overfitting -> the ML model captures the noise in your training dataset
- The **regularization** term is an addition to the loss function which helps generalize the model
 - **L1** or Lasso regularization adds a penalty which is the sum of the absolute values of the weights
 - L1+MSE
$$\text{Min}(\sum_{i=1}^n (y_i - w_i x_i)^2 + p \sum_{i=1}^n |w_i|)$$
 - **L2** or Ridge regularization adds a penalty which is the sum of the squared values of weights
 - L2+MSE
$$\text{Min}(\sum_{i=1}^n (y_i - w_i x_i)^2 + p \sum_{i=1}^n (w_i)^2)$$
- **Dropout** in NN context: hidden nodes are dropped randomly
- **Early Stopping** is a time regularization technique which stops training based on given criteria

Activation functions

- Classification: sigmoid functions
 - sigmoids and tanh functions are sometimes avoided due to the vanishing gradient problem
- **ReLU** function is a general activation function
- dead neurons in our networks -> the leaky ReLU
- ReLU function should only be used in the hidden layers
- As a rule of thumb, start with ReLU


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$



Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

